


# Floyd's algorithm

# Overview


- Chapter 6 from *Michael J. Quinn, Parallel Programming in C with MPI and OpenMP*



- Floyd's algorithm: solving *the all-pairs shortest-path problem*

# Finding shortest paths

- Starting point: a graph of vertices and weighted edges



- Each edge is of a direction and has a length
  - if there's path from vertex  $i$  to  $j$ , there may not be path from vertex  $j$  to  $i$
  - path length from vertex  $i$  to  $j$  may be different than path length from vertex  $j$  to  $i$
- Objective: finding the shortest path between every pair of vertices ( $i \rightarrow j$ )
- Application: table of driving distances between city pairs

# Adjacency matrix

- There are  $n$  vertices
- The direct path length from vertex  $i$  to vertex  $j$  is stored as  $a[i, j]$
- An  $n \times n$  adjacency matrix  $a$  keeps the entire connectivity info

	0	1	2	3	4	5
0	0	2	5	$\infty$	$\infty$	$\infty$
1	$\infty$	0	7	1	$\infty$	8
2	$\infty$	$\infty$	0	4	$\infty$	$\infty$
3	$\infty$	$\infty$	$\infty$	0	3	$\infty$
4	$\infty$	$\infty$	2	$\infty$	0	3
5	$\infty$	5	$\infty$	2	4	0

- If  $a[i, j]$  is  $\infty$ , it means there is no direct path from vertex  $i$  to vertex  $j$

## Example of all-pairs shortest path

For the adjacency matrix given on the previous slide, the solution of the all-pairs shortest path is as follows:

	0	1	2	3	4	5
0	0	2	5	3	6	9
1	$\infty$	0	6	1	4	7
2	$\infty$	15	0	4	7	10
3	$\infty$	11	5	0	3	6
4	$\infty$	8	2	5	0	3
5	$\infty$	5	6	2	4	0

Table of shortest path lengths

# Floyd's algorithm

Input:  $n$  — number of vertices

$a$  — adjacency matrix

Output: Transformed  $a$  that contains the shortest path lengths

```
for  $k \leftarrow 0$  to  $n - 1$ 
  for  $i \leftarrow 0$  to  $n - 1$ 
    for  $j \leftarrow 0$  to  $n - 1$ 
       $a[i, j] \leftarrow \min(a[i, j], a[i, k] + a[k, j])$ 
    endfor
  endfor
endfor
```

# Some observations

- Floyd's algorithm is an exhaustive and incremental approach
- The entries of the  $a$ -matrix are updated  $n$  rounds
- $a[i, j]$  is compared with all  $n$  possibilities,  
that is, against  $a[i, k] + a[k, j]$ , for  $0 \leq k \leq n - 1$
- $n^3$  of comparisons in total

# Source of parallelism

- During the  $k$ 'th iteration, the work is (in C syntax)

```
for (i=0; i<n; i++)  
    for (j=0; j<n, j++)  
        a[i][j] = MIN( a[i][j], a[i][k]+a[k][j] );
```


- Can all the entries in  $a$  be updated concurrently?
- Yes, because the  $k$ 'th column and the  $k$ 'th row remain the same during the  $k$ 'th iteration!
  - Note that  $a[i][k] = \text{MIN}(a[i][k], a[i][k] + a[k][j])$  will be the same as  $a[i][k]$
  - Note that  $a[k][j] = \text{MIN}(a[k][j], a[k][k] + a[k][j])$  will be the same as  $a[k][j]$



# Design of a parallel algorithm

Using Foster's design methodology:

- Partitioning — each  $a[i, j]$  is a primitive task
- Communication — during the  $k$ 'th iteration, updating  $a[i, j]$  needs values of  $a[i, k]$  and  $a[k, j]$ 
  - broadcast  $a[k, j]$  to  $a[0, j], a[1, j], \dots, a[n-1, j]$
  - broadcast  $a[i, k]$  to  $a[i, 0], a[i, 1], \dots, a[i, n-1]$



# Agglomeration and mapping

- Let one MPI process be responsible for a piece of the  $a$  matrix
- Memory storage of  $a$  is accordingly divided
- The division can in principle be arbitrary, as long as the number of all  $a[i, j]$  entries is divided evenly
- However, a row-wise block data division is very convenient
  - 2D arrays in C are row-major
  - easy to send/receive an entire row of  $a$
- We therefore choose to assign one MPI process with a number of consecutive rows of  $a$

# Communication pattern

- Recall that in the  $k$ 'th iteration:
$$a[i, j] \leftarrow \min(a[i, j], a[i, k] + a[k, j])$$
- Since entries of  $a$  are divided into rowwise blocks, so  $a[i, k]$  is also in the memory of the MPI process that owns  $a[i, j]$
- However,  $a[k, j]$  is probably in another MPI process's memory
- Communication is therefore needed!
  - Before the  $k$ 'th iteration, the MPI process that owns the  $k$ 'th row of the  $a$  matrix should broadcast this row to everyone else

# Recap: creating 2D arrays in C

To create a 2D array with  $m$  rows and  $n$  columns:

```
int **B, *Bstorage, i;  
...  
Bstorage=(int*)malloc(m*n*sizeof(int));  
B=(int**)malloc(m*sizeof(int*));  
for (i=0; i<m; i++)  
    B[i] = &Bstorage[i*n];
```

The underlying storage is contiguous, making it possible to send and receive an entire 2D array.

# Global index vs. local index

- Suppose a matrix (2D array) is divided into row-wise blocks and distributed among  $p$  MPI processes
- Process  $i$  only allocates storage for its assigned row block
  - from row  $\lfloor (i \cdot n)/p \rfloor$  of matrix  $a$  until row  $\lfloor ((i + 1) \cdot n)/p \rfloor - 1$
- We need to know: Which global row does a local row correspond to?
- Mapping: local index  $\rightarrow$  global index
- On process number `proc_id`  
`global_index=BLOCK_LOW(proc_id, p, n)+local_index`

# Main work of parallel Floyd's algorithm

```
void compute_shortest_paths (int id, int p, dtype **a, int n)
{
    int i, j, k;
    int offset; /* Local index of broadcast row */
    int root; /* Process controlling row to be bcast */
    int* tmp; /* Holds the broadcast row */
    tmp = (dtype *) malloc (n * sizeof(dtype));
    for (k = 0; k < n; k++) {
        root = BLOCK_OWNER(k,p,n);
        if (root == id) {
            offset = k - BLOCK_LOW(id,p,n);
            for (j = 0; j < n; j++)
                tmp[j] = a[offset][j];
        }
        MPI_Bcast (tmp, n, MPI_TYPE, root, MPI_COMM_WORLD);
        for (i = 0; i < BLOCK_SIZE(id,p,n); i++)
            for (j = 0; j < n; j++)
                a[i][j] = MIN(a[i][j],a[i][k]+tmp[j]);
    }
    free (tmp);
}
```

# Matrix input

- Recall that each MPI process only stores a part of the  $a$  matrix
- When reading  $a$  from a file, we can
  - let only process  $p - 1$  do the input
  - once the number of rows needed by process  $i$  are read in, they are sent from process  $p - 1$  to process  $i$  using `MPI_Send`
  - process  $i$  must issue a matching `MPI_Recv`
- The above simple strategy is not parallel
- Parallel I/O can be done using MPI-2 commands

# Matrix output

- For example, we let only process 0 do the output
- Each process needs to send its part of  $a$  to process 0
- To avoid many processes sending its entire subdata to process 0 at the same time
  - Process 0 communicates with the other processes in turn
  - Each process waits for a “hint” (a short message) from process 0 before sending its data (a large message)



# Deadlock

## ● Typical deadlock example 1

```
if (rank==0) {
    MPI_Recv(&b,1,MPI_INT,1,tag_b,MPI_COMM_WORLD,&status)
    MPI_Send(&a,1,MPI_INT,1,tag_a,MPI_COMM_WORLD);
} else if (rank==1) {
    MPI_Recv(&a,1,MPI_INT,0,tag_a,MPI_COMM_WORLD,&status)
    MPI_Send(&b,1,MPI_INT,0,tag_b,MPI_COMM_WORLD);
}
```

## ● Typical deadlock example 2

```
if (rank==0) {
    MPI_Send(&a,1,MPI_INT,1,1,MPI_COMM_WORLD);
    MPI_Recv(&b,1,MPI_INT,1,1,MPI_COMM_WORLD,&status);
} else if (rank==1) {
    MPI_Send(&b,1,MPI_INT,0,0,MPI_COMM_WORLD);
    MPI_Recv(&a,1,MPI_INT,0,0,MPI_COMM_WORLD,&status);
}
```

# Analysis

- Serial algorithm time usage:  $n^3\chi$
- Parallel algorithm
  - non-communication time usage:  $n^2\lceil n/p\rceil\chi$
  - communication (broadcast) time usage:  $n\lceil\log_2 p\rceil(\lambda + 4n/\beta)$ 
    - assuming each entry of matrix  $a$  needs 4 bytes
    - assuming  $\lambda$  as communication latency
    - assuming  $\beta$  as communication bandwidth (# bytes per second)
- Read Section 6.7 for a more detailed analysis that allows overlap between computation and communication

# Exercises

- Write an MPI program that uses  $p$  processes to produce a JPEG picture of  $n \times n$  pixels. The picture should have white background and a black circle (of radius  $n/4$ ) in the middle. (The existing C code collection <http://heim.ifi.uio.no/xingca/inf-verk3830/simple-jpeg.tar.gz> can be used.)
- Implement the complete Floyd's algorithm and try it on a large enough adjacency matrix.